

Chapter 4: Threads & Concurrency





Chapter 4: Threads

- Overview
- Multicore Programming
- Multithreading Models
- Threading Issues
- Operating System Examples





Objectives

- To introduce the notion of a thread—a fundamental unit of CPU utilization that forms the basis of multithreaded computer systems
- To examine issues related to multithreaded programming





Motivation

- ❑ Many modern applications are multithreaded
- ❑ Threads run within application
- ❑ Multiple tasks within an application (process) can be implemented by separate threads
 - ❑ Update display
 - ❑ Fetch data
 - ❑ Spell checking
 - ❑ Answer a network request
- ❑ Process creation is heavy-weight while thread creation is light-weight
- ❑ Can simplify code, increase efficiency
- ❑ Kernels are generally multithreaded





Benefits


- ❑ **Responsiveness** – may allow continued execution if part of process is blocked, especially important for user interfaces
- ❑ **Resource Sharing** – threads share resources of process, easier than shared memory or message passing
- ❑ **Economy** – cheaper than process creation, thread switching lower overhead than context switching
- ❑ **Scalability** – process can take advantage of multicore architectures





A Multi-Core Design


- ❑ Multi-chip and **multicore**
 - ❑ On-chip communications faster
 - ❑ Power consumption can be reduced





A Modern Processor

□ Intel Sandy Bridge



How to program this?





Multicore Programming


- ❑ **Multicore** or **multiprocessor** systems, programming challenges include:
 - ❑ Dividing activities – how to divide tasks
 - ❑ Balance – each task perform “equal” amount of work
 - ❑ Data splitting
 - ❑ Data dependency – synchronization
 - ❑ Testing and debugging – different path of executions
- ❑ **Parallelism** a system can perform more than one task simultaneously
 - ❑ **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
 - ❑ **Task parallelism** – distributing threads across cores, each thread performing unique operation
- ❑ **Concurrency** supports more than one task making progress
 - ❑ Multiplex over the time
 - ❑ Single processor / core, scheduler providing concurrency






Concurrency vs. Parallelism

- Concurrent execution on single-core system:





- Parallelism on a multi-core system:





Data and Task Parallelism

- **Data parallelism** – distributes subsets of the same data across multiple cores, same operation on each
- **Task parallelism** – distributing threads across cores, each thread performing unique operation





Amdahl's Law

- Identifies performance gains from adding additional cores to an application that has both serial and parallel components
- S is serial portion
- N processing cores

$$speedup \leq \frac{1}{S + \frac{(1-S)}{N}}$$

- That is, if application is 75% parallel / 25% serial, moving from 1 to 2 cores results in speedup of 1.6 times
- As N approaches infinity, speedup approaches $1 / S$


Serial portion of an application has disproportionate effect on performance gained by adding additional cores

- But does the law take into account contemporary multicore systems?



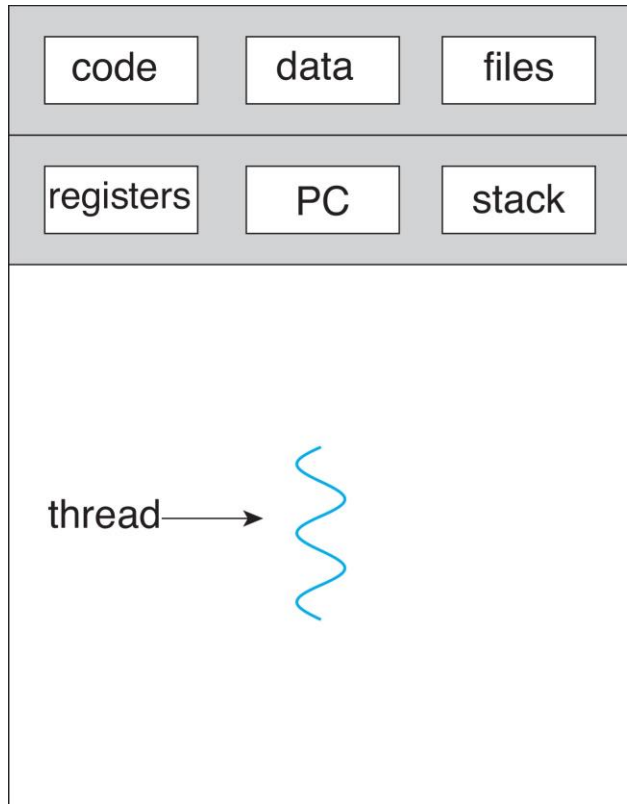


Amdahl's Law

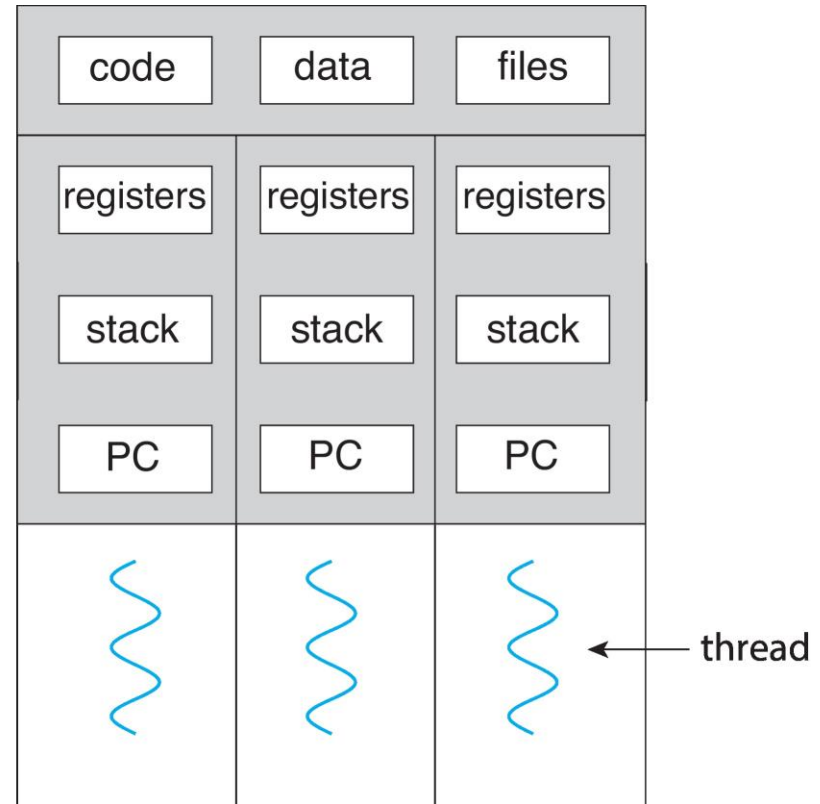




Single and Multithreaded Processes



single-threaded process



multithreaded process





Thread

- ❑ **Thread**: Single unique execution context (“lightweight process”)
 - ❑ Program Counter, Registers, Execution Flags, Stack
 - ❑ A thread is executing on a processor when it is resident in the registers.
 - ❑ PC register holds the address of executing instruction in the thread
 - ❑ Registers hold the root state of the thread (other state in memory)

- ❑ Each Thread has a **Thread Control Block** (TCB)
 - ❑ **Execution state**: CPU registers, program counter, pointer to stack
 - ❑ **Scheduling info**: State (more later), priority, CPU time
 - ❑ **Accounting Info**
 - ❑ **Various Pointers** (for implementing scheduling queues)
 - ❑ **Pointer to enclosing process**: PCB

- ❑ In Nachos: “thread” is a class that includes the TCB

- ❑ OS keeps track of TCBs in protected memory region
 - ❑ Array, or Linked List, or ...





Thread State

- ❑ Threads encapsulate **concurrency**: “**Active**” component
- ❑ Address spaces encapsulate protection: “**Passive**” part
 - ❑ Keeps buggy program from thrashing the system
- ❑ State shared by all threads in process/address space
 - ❑ Contents of memory (global variables, heap)
 - ❑ I/O state (file descriptors, network connections, etc.)
- ❑ State “private” to each thread
 - ❑ Kept in **TCB** \equiv **Thread Control Block**
 - ❑ CPU registers (including, program counter)
 - ❑ Execution stack (parameters, temporary variables, PC saved)





Shared vs. Per-Thread State

Shared State

Heap

Global Variables

Code

Per-Thread State

Thread Control Block (TCB)

Stack Information

Saved Registers

Thread Metadata

Stack

Per-Thread State

Thread Control Block (TCB)

Stack Information

Saved Registers

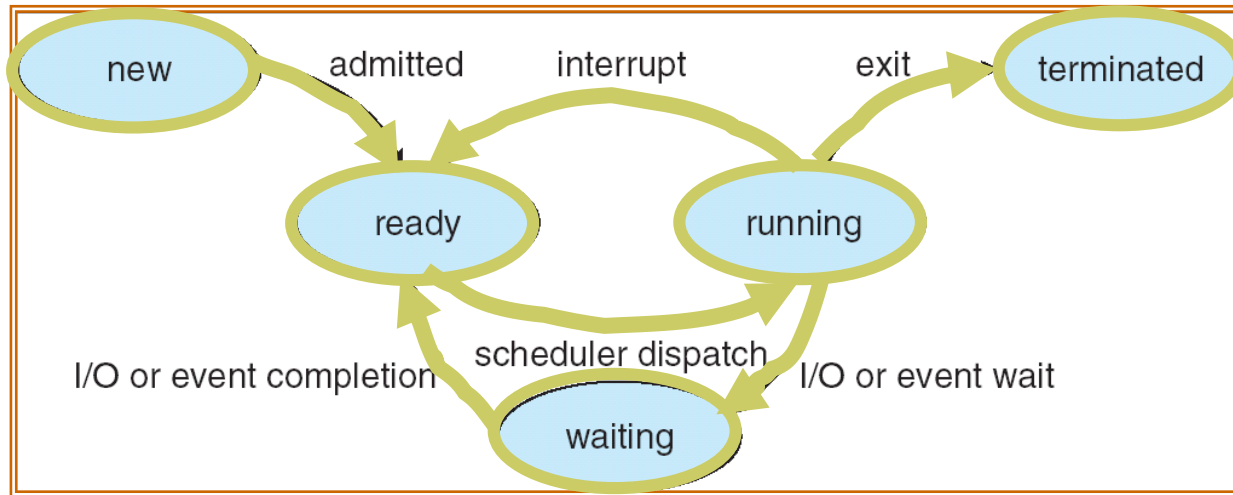
Thread Metadata

Stack





Lifecycle of a Thread



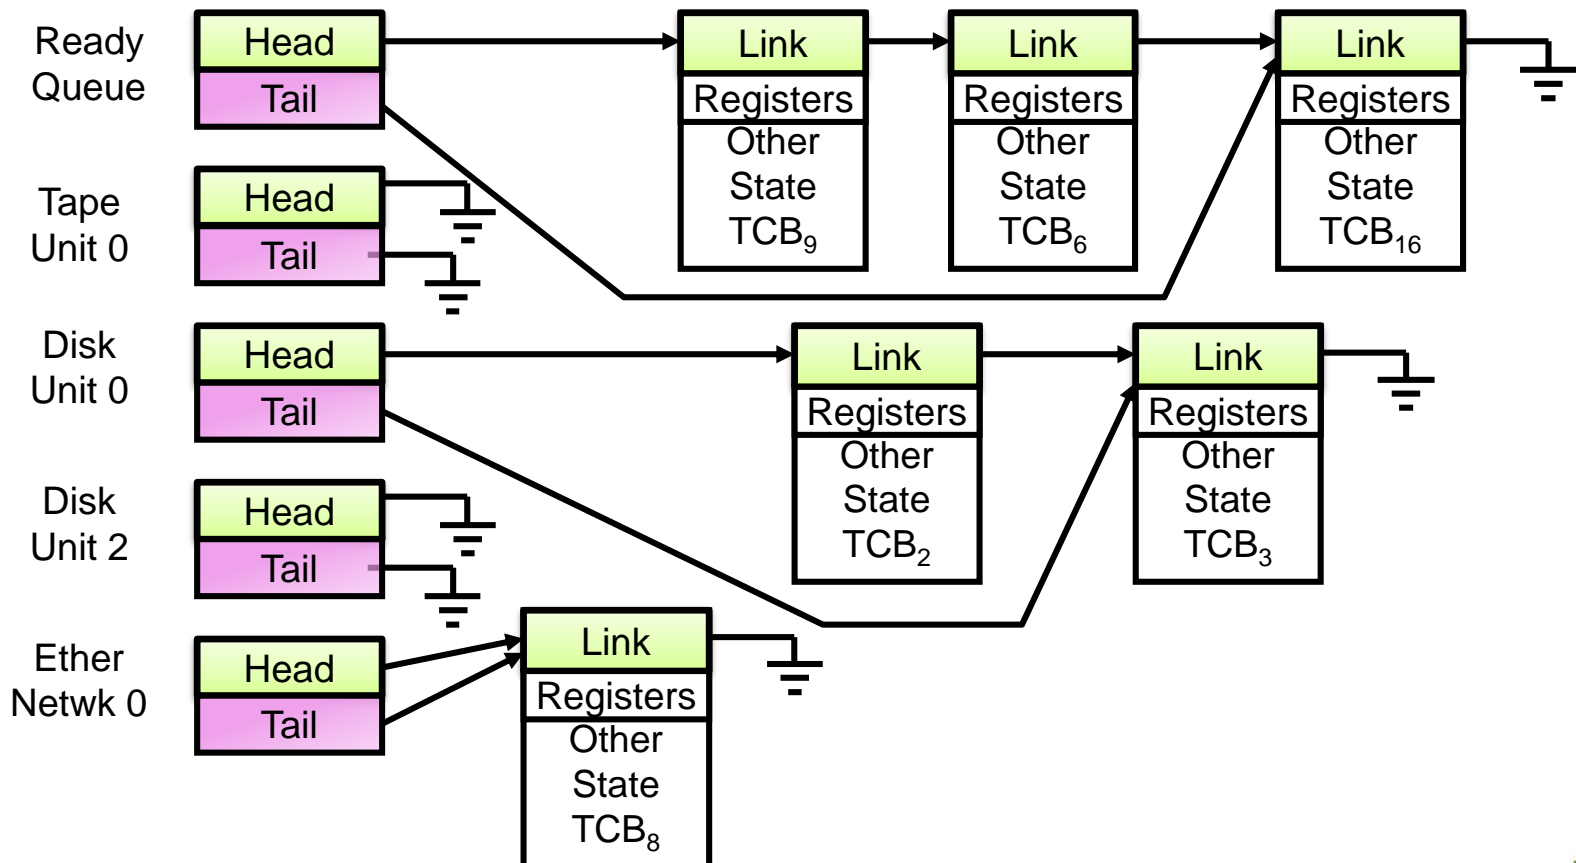
- As a thread executes, it changes state:
 - **new**: The thread is being created
 - **ready**: The thread is waiting to run
 - **running**: Instructions are being executed
 - **waiting**: Thread waiting for some event to occur
 - **terminated**: The thread has finished execution
- “Active” threads are represented by their TCBs
 - TCBs organized into queues based on their states






Ready Queue And Various I/O Device Queues

- Thread not running \Rightarrow TCB is in some scheduler queue
 - Separate queue for each device/signal/condition
 - Each queue can have a different scheduler policy





Examples of Multithreaded Programs


- ❑ Embedded systems
 - ❑ Elevators, Planes, Medical systems, Wristwatches
 - ❑ Single Program, concurrent operations
- ❑ Most modern OS kernels
 - ❑ Internally concurrent to deal with concurrent requests by multiple users
 - ❑ But no protection needed within kernel
- ❑ Database Servers
 - ❑ Access to shared data by many concurrent users
 - ❑ Also background utility processing must be done
- ❑ Network Servers
 - ❑ Concurrent requests from network
 - ❑ Again, single program, multiple concurrent operations
 - ❑ File server, Web server, and airline reservation systems
- ❑ Parallel Programming (more than one physical CPU)
 - ❑ Split a program into multiple threads for parallelism





User Threads and Kernel Threads

- **User threads** - management done by user-level threads library
- Three primary thread libraries:
 - POSIX **Pthreads**, Windows threads and Java threads
- **Kernel threads** - supported by the operating system
- Examples – virtually all general purpose operating systems, including:
 - Windows, Linux, Mac OS X iOS, Android





Multithreading Models


- Many-to-One
- One-to-One
- Many-to-Many





Many-to-One


- ❑ Many user-level threads mapped to single kernel thread - **scheduling**
- ❑ One thread blocking causes all to block
- ❑ Multiple threads may not run in parallel on muticore system because only one can be in kernel at a time
- ❑ Few systems currently use this model
- ❑ Examples:
 - ❑ Solaris Green Threads
 - ❑ GNU Portable Threads





One-to-One


- ❑ Each user-level thread maps to kernel thread
- ❑ Creating a user-level thread creates a kernel thread
- ❑ More concurrency than many-to-one
- ❑ Number of threads per process sometimes restricted due to overhead
- ❑ Examples
 - ❑ Windows
 - ❑ Linux





Many-to-Many Model


- Allows many user level threads to be mapped to many kernel threads
M:N
- Allows the operating system to create a sufficient number of kernel threads
- Windows with the *ThreadFiber* package
- Otherwise not very common





Two-level Model

- Similar to **M:N**, except that it allows a user thread to be **bound** to kernel thread





Threading Issues

- ❑ Semantics of `fork()` and `exec()` system calls
- ❑ Signal handling
 - ❑ Synchronous and asynchronous
- ❑ Thread cancellation of target thread
 - ❑ Asynchronous or deferred
- ❑ Thread-local storage
- ❑ Scheduler Activations





Semantics of `fork()` and `exec()`

- Does `fork()` duplicate only the calling thread or all threads?
 - Some UNIXes have two versions of `fork()`
- `exec()` usually works as normal – replace the running process including all threads





Signal Handling

- n **Signals** are used in UNIX systems to notify a process that a particular event has occurred.
- n A **signal handler** is used to process signals
 1. Signal is generated by a particular event (e.g., process termination)
 2. Signal is delivered to a process
 3. Signal is handled by one of two signal handlers:
 1. default
 2. user-defined
- n Every signal has **default handler** that kernel runs when handling signal
 - | **User-defined signal handler** can override default
 - | For single-threaded, signal delivered to process





Signal Handling (Cont.)

- n Where should a signal be delivered for multi-threaded?
 - | Deliver the signal to the thread to which the signal applies
 - | Deliver the signal to every thread in the process
 - | Deliver the signal to certain threads in the process
 - | Assign a specific thread to receive all signals for the process





Thread Cancellation

- ❑ Terminating a thread before it has finished
- ❑ Thread to be canceled is **target thread**
- ❑ Two general approaches:
 - ❑ **Asynchronous cancellation** terminates the target thread immediately
 - ❑ **Deferred cancellation** allows the target thread to periodically check if it should be cancelled
- ❑ Pthread code to create and cancel a thread:

```
pthread_t tid;

/* create the thread */
pthread_create(&tid, 0, worker, NULL);

. . .

/* cancel the thread */
pthread_cancel(tid);

/* wait for the thread to terminate */
pthread_join(tid, NULL);
```





Thread Cancellation (Cont.)

- pThread code to create and cancel a thread:
 - pthreads: POSIX standard for thread programming
 - Need to #include <pthread.h>

```
pthread_t tid;  
  
/* create the thread */  
pthread_create(&tid, 0, worker, NULL);  
  
. . .  
  
/* cancel the thread */  
pthread_cancel(tid);  
  
/* wait for the thread to terminate */  
pthread_join(tid, NULL);
```

Start routine
↓





Thread Cancellation (Cont.)

- Invoking thread cancellation requests cancellation, but actual cancellation depends on thread **state** and **type**

| Mode | State | Type |
|--------------|----------|--------------|
| Off | Disabled | – |
| Deferred | Enabled | Deferred |
| Asynchronous | Enabled | Asynchronous |

- If thread has cancellation disabled, cancellation remains pending until thread enables it
- Default type is *deferred*
 - Cancellation only occurs when thread reaches **cancellation point**
 - ▶ I.e., `pthread_testcancel()`
 - ▶ Then **cleanup handler** is invoked
- On Linux systems, thread cancellation is handled through signals





Thread-Local Storage


- ❑ Thread-local storage (TLS) allows each thread to have its own copy of data
- ❑ Useful when you do not have control over the thread creation process (i.e., when using a thread pool)
- ❑ Different from local variables
 - ❑ Local variables visible only during single function invocation
 - ❑ TLS visible across function invocations
- ❑ Similar to **static** data
 - ❑ TLS is unique to each thread





Scheduler Activations

- ❑ Both M:N and Two-level models require communication to maintain the appropriate number of kernel threads allocated to the application
- ❑ Typically use an intermediate data structure between user and kernel threads – **lightweight process (LWP)**
 - ❑ Appears to be a virtual processor on which process can schedule user thread to run
 - ❑ Each LWP attached to kernel thread
 - ❑ How many LWPs to create?
- ❑ Scheduler activations provide **upcalls** - a communication mechanism from the kernel to the **upcall handler** in the thread library
- ❑ This communication allows an application to maintain the correct number kernel threads





Operating System Examples

- Windows Threads
- Linux Threads





Windows Threads

- ❑ Windows API – primary API for Windows applications
- ❑ Implements the one-to-one mapping, kernel-level
- ❑ Each thread contains
 - ❑ A thread id
 - ❑ Register set representing state of processor
 - ❑ Separate user and kernel stacks for when thread runs in user mode or kernel mode
 - ❑ Private data storage area used by run-time libraries and dynamic link libraries (DLLs)
- ❑ The register set, stacks, and private storage area are known as the **context** of the thread





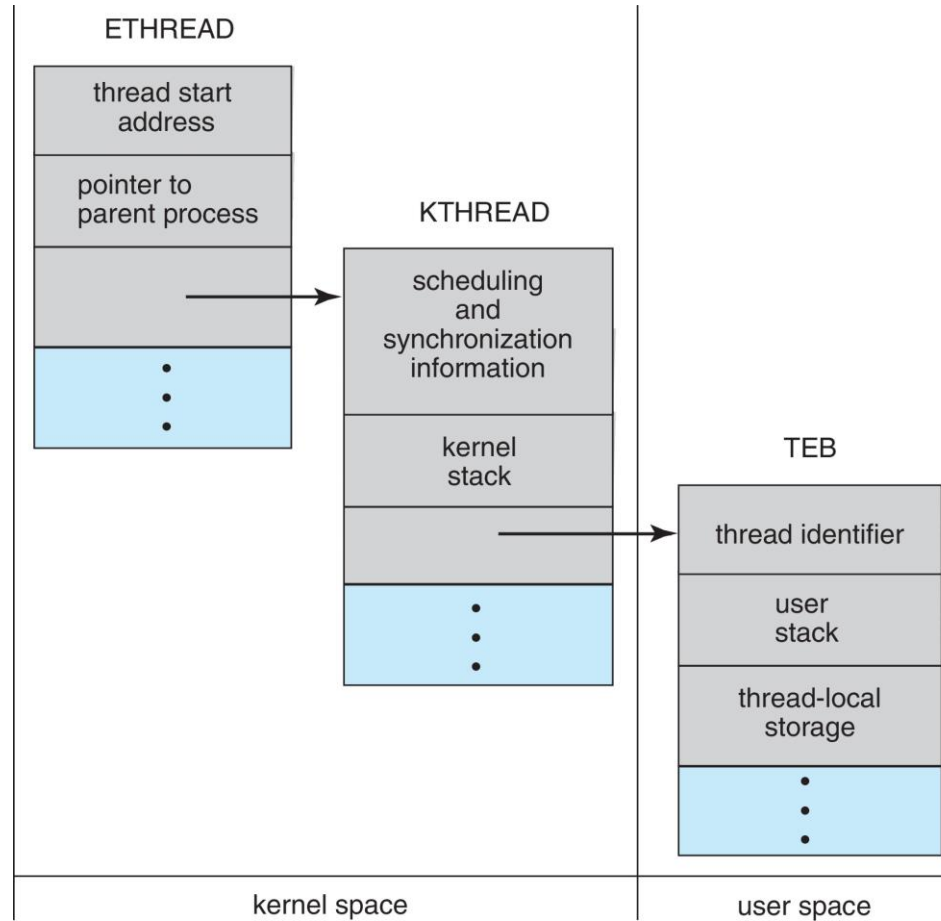
Windows Threads (Cont.)

- The primary data structures of a thread include:
 - ETHREAD (executive thread block) – includes pointer to process to which thread belongs and to KTHREAD, in kernel space
 - KTHREAD (kernel thread block) – scheduling and synchronization info, kernel-mode stack, pointer to TEB, in kernel space
 - TEB (thread environment block) – thread id, user-mode stack, thread-local storage, in user space





Windows Threads Data Structures





Linux Threads

- ❑ Linux refers to them as *tasks* rather than *threads*
- ❑ Thread creation is done through `clone()` system call
- ❑ `clone()` allows a child task to share the part of the execution context with the parent, such as address space, file descriptors, signal handler
- ❑ Flags control behavior

| flag | meaning |
|---------------|------------------------------------|
| CLONE_FS | File-system information is shared. |
| CLONE_VM | The same memory space is shared. |
| CLONE_SIGHAND | Signal handlers are shared. |
| CLONE_FILES | The set of open files is shared. |

```
int clone(int (*fn)(void *), void *child_stack,  
          int flags, void *arg, ... );
```

- ❑ The child thread executes the function `fn`, `arg` specified the arguments for the function `fn`.



End of Chapter 4

