

Chapter 5: CPU Scheduling





Chapter 5: CPU Scheduling

- ❑ Basic Concepts
- ❑ Scheduling Criteria
- ❑ Scheduling Algorithms
- ❑ Thread Scheduling
- ❑ Multi-Processor Scheduling
- ❑ Real-Time CPU Scheduling
- ❑ Operating Systems Examples
- ❑ Algorithm Evaluation





Objectives

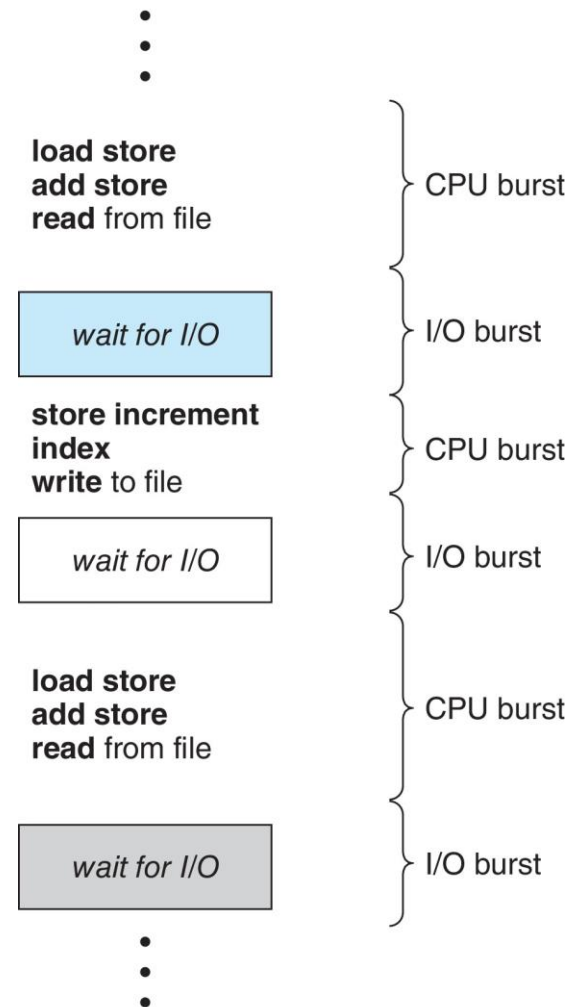
- ❑ Describe various CPU scheduling algorithms
- ❑ Assess CPU scheduling algorithms based on scheduling criteria
- ❑ Explain the issues related to multiprocessor and multicore scheduling
- ❑ Describe various real-time scheduling algorithms
- ❑ Apply modeling and simulations to evaluate CPU scheduling algorithms





Basic Concepts

- ❑ Maximum CPU utilization obtained with multiprogramming
- ❑ CPU–I/O Burst Cycle – Process execution consists of a **cycle** of CPU execution and I/O wait
- ❑ **CPU burst** followed by **I/O burst**
- ❑ CPU burst distribution is of primary concern






Histogram of CPU-burst Times

Large number of short bursts

Small number of longer bursts





CPU Scheduler


- The **CPU scheduler** selects from among the processes in ready queue, and allocates the a CPU core to one of them
 - Queue may be ordered in various ways
- CPU scheduling decisions may take place when a process:
 1. Switches from running to waiting state
 2. Switches from running to ready state
 3. Switches from waiting to ready
 4. Terminates
- Scheduling under 1 and 4 is **nonpreemptive**
- All other scheduling is **preemptive**
 - Consider access to shared data
 - Consider preemption while in kernel mode
 - Consider interrupts occurring during crucial OS activities





Dispatcher

- ❑ Dispatcher module gives control of the CPU to the process selected by the short-term scheduler; this involves:
 - ❑ switching context
 - ❑ switching to user mode
 - ❑ jumping to the proper location in the user program to restart that program
- ❑ **Dispatch latency** – time it takes for the dispatcher to stop one process and start another running





Scheduling Criteria

- **CPU utilization** – keep the CPU as busy as possible
- **Throughput** – # of processes or jobs completed per time unit (second)
 - Efficient use of the resources (CPU, memory, disk, etc.)
 - Minimize the overhead (for instance, context switching)
- **Waiting time** – amount of time a process waiting in the ready queue
- **Turnaround time** – amount of time to execute a particular process, usually measured by the CPU burst time plus waiting time
 - For single CPU burst, turnaround time = waiting time + CPU burst time
- **Response time** – amount of time it takes from when a request was submitted until the **first** response is produced
 - Time to echo a keystroke in editor
 - Time to compile a program
- **Fairness**
 - Resources such as CPU are utilized in some “fair” manner






First-Come, First-Served (FCFS) Scheduling

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

- Suppose that the processes arrive in the order: P_1 , P_2 , P_3
The Gantt Chart for the schedule is:



- Waiting time for $P_1 = 0$; $P_2 = 24$; $P_3 = 27$
- Average waiting time: $(0 + 24 + 27)/3 = 17$
- In earlier systems, FCFS means that one program is scheduled to run until completion (including all I/O)
- Now this means a process finishes its current CPU burst time





FCFS Scheduling (Cont.)

Suppose that the processes arrive in the order:

$$P_2, P_3, P_1$$

□ The Gantt chart for the schedule is:



- Waiting time for $P_1 = 6$; $P_2 = 0$; $P_3 = 3$
- Average waiting time: $(6 + 0 + 3)/3 = 3$
- Much better than previous case
- **Convoy effect** - short process stuck behind long process, potentially bad for short jobs!
 - Consider one CPU-bound and many I/O-bound processes
 - Waiting in banks: depositing a check, stuck behind new account opening





Round Robin (RR)

- ❑ Each process gets a small unit of CPU time (**time quantum q**), usually 10-100 milliseconds. After this time has elapsed, the process is preempted and added to the end of the ready queue.
- ❑ Given n processes, each process gets $1/n$ of the CPU time in chunks of at most q time units at once.
 - ❑ No process waits more than $(n-1)q$ time units.
- ❑ Timer interrupts every quantum to schedule next process, or the process blocks upon completing its current CPU burst time ($< q$)
- ❑ Performance
 - ❑ q large \Rightarrow FIFO
 - ❑ q small \Rightarrow interleaved, but q must be large with respect to context switch, otherwise overhead is too high






Example of RR with Time Quantum = 4

<u>Process</u>	<u>Burst Time</u>
P_1	24
P_2	3
P_3	3

□ The Gantt chart is:




- Waiting time for $P_1=6$, $P_2=4$, $P_3=7$
- Average waiting time $(6+4+7)/3 = 5.67$
- **Response time** for $P_1=4$, $P_2=7$, $P_3=10$, average = 7
- q should be large compared to context switch time
- q usually 10ms to 100ms, context switch < 10 usec
- Better for short job (fair), context switching adds up for long jobs





Time Quantum and Context Switch Time





Comparisons b/t FCFS and RR

Assuming zero-cost context-switching time, is RR **always better** than FCFS?

- **Simple example:** 10 jobs starting at the same time, each taking 100s of CPU time; RR scheduler quantum of 1s;


Job #	FIFO	RR
1	100	991
2	200	992
...
9	900	999
10	1000	1000

- **Average job completion time is much worse under RR!**
 - ▶ Bad when all jobs have the same length





Earlier Example w/ Different Quantum




	Quantum	P ₁	P ₂	P ₃	P ₄	Average
Wait Time	Best FCFS	32	0	85	8	31¼
	Q = 1	84	22	85	57	62
	Q = 5	82	20	85	58	61¼
	Q = 8	80	8	85	56	57¼
	Q = 10	82	10	85	68	61¼
	Q = 20	72	20	85	88	66¼
	Worst FCFS	68	145	0	121	83½
Completion Time	Best FCFS	85	8	153	32	69½
	Q = 1	137	30	153	81	100½
	Q = 5	135	28	153	82	99½
	Q = 8	133	16	153	80	95½
	Q = 10	135	18	153	92	99½
	Q = 20	125	28	153	112	104½
	Worst FCFS	121	153	68	145	121¾





Turnaround Time Varies With The Time Quantum



process	time
P_1	6
P_2	3
P_3	1
P_4	7

- The average turnaround time does not necessarily improve as the time quantum size increases
- In general, the average turnaround time can be improved if most processes finish their current CPU bursts within a single quantum
- The time quantum can not be too big, in which RR degenerates to an FCFS policy
- **A rule of thumb:** 80% of CPU bursts should be shorter than the time quantum q





Shortest-Job-First (SJF) Scheduling

- What if we knew the future
- Associate with each process the length of its next CPU burst
 - To schedule the process with the shortest CPU burst
- SJF is **optimal** – yields the minimum average waiting time for a given set of processes
 - The difficulty is knowing the length of the next CPU request
 - Could ask the user
 - The basic idea is to get the short jobs out of system sooner
 - Big effect on short jobs, only small effect on long jobs
 - This can be applied to whole program or current CPU burst






Example of SJF

<u>Process</u>	<u>Burst Time</u>
P_1	6
P_2	8
P_3	7
P_4	3

□ SJF scheduling chart



- Average waiting time = $(3 + 16 + 9 + 0) / 4 = 7$
- The (best) FCFS can perform the same if arrival order is the same






Determining Length of Next CPU Burst

- Can estimate the length based on the past behavior
 - Then pick process with shortest predicted next CPU burst
- Can be done by using the length of previous CPU bursts, and exponential averaging algorithm
 1. t_n = actual length of n^{th} CPU burst
 2. τ_{n+1} = predicted value for the next CPU burst
 3. $\alpha, 0 \leq \alpha \leq 1$
 4. Define : $\tau_{n+1} = \alpha t_n + (1 - \alpha)\tau_n$.
- Commonly, α set to $\frac{1}{2}$
- Preemptive version called **shortest-remaining-time-first (SRTF)**





Prediction of the Length of the Next CPU Burst



CPU burst (t_i)	6	4	6	4	13	13	13	...	
"guess" (τ_i)	10	8	6	6	5	9	11	12	...





Examples of Exponential Averaging

- $\alpha = 0$
 - $\tau_{n+1} = \tau_n$
 - Recent history does not count
- $\alpha = 1$
 - $\tau_{n+1} = \alpha t_n$
 - Only the actual last CPU burst counts
- If we expand the formula, we get:

$$\begin{aligned}\tau_{n+1} = & \alpha t_n + (1 - \alpha)\alpha t_{n-1} + \dots \\ & + (1 - \alpha)^j \alpha t_{n-j} + \dots \\ & + (1 - \alpha)^{n+1} \tau_0\end{aligned}$$

- Since both α and $(1 - \alpha)$ are less than or equal to 1, each successive term has less weight than its predecessor, thus its effect is diminishing exponentially fast






Example of Shortest-remaining-time-first

- Now we add the concepts of varying arrival times and preemption to the analysis

<u>Process</u>	<u>Arrival Time</u>	<u>Burst Time</u>
P_1	0	8
P_2	1	4
P_3	2	9
P_4	3	5

- Preemptive SJF Gantt Chart*



- Average waiting time = $[(10-1)+(1-1)+(17-2)+5-3]/4 = 26/4 = 6.5$ msec
- Now scheduling needs to be considered with each arrival





Comparison of SJF/SRTF and FCFS

- ❑ SJF/SRTF are the best we can do at minimizing the average waiting time (or the average turnaround time)
 - ❑ Provably optimal (SJF among non-preemptive, SRTF among preemptive)
 - ❑ SRTF is always at least as good as SJF
- ❑ SJF/SRTF performs the same as FCFS if all processes have the same CPU burst times
- ❑ SJF/SRTF can lead to starvation





Priority Scheduling

- A priority number (integer) is associated with each process
- The CPU is allocated to the process with the highest priority (smallest integer \equiv highest priority), it can be
 - Preemptive (upon new arrival from higher priority)
 - Nonpreemptive
- SJF is priority scheduling where priority is the inverse of predicted next CPU burst time
- Problem \equiv **Starvation** – low priority processes may never execute
- Solution \equiv **Aging** – as time progresses increase the priority of the process






Example of Priority Scheduling

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	10	3
P_2	1	1
P_3	2	4
P_4	1	5
P_5	5	2

□ Priority scheduling Gantt Chart



□ Average waiting time = 8.2 msec






Priority Scheduling w/ Round-Robin

<u>Process</u>	<u>Burst Time</u>	<u>Priority</u>
P_1	4	3
P_2	5	2
P_3	8	2
P_4	7	1
P_5	3	3


- ❑ Run the process with the highest priority. Processes with the same priority run round-robin
- ❑ Gantt Chart with 2 ms time quantum





Multilevel Queue

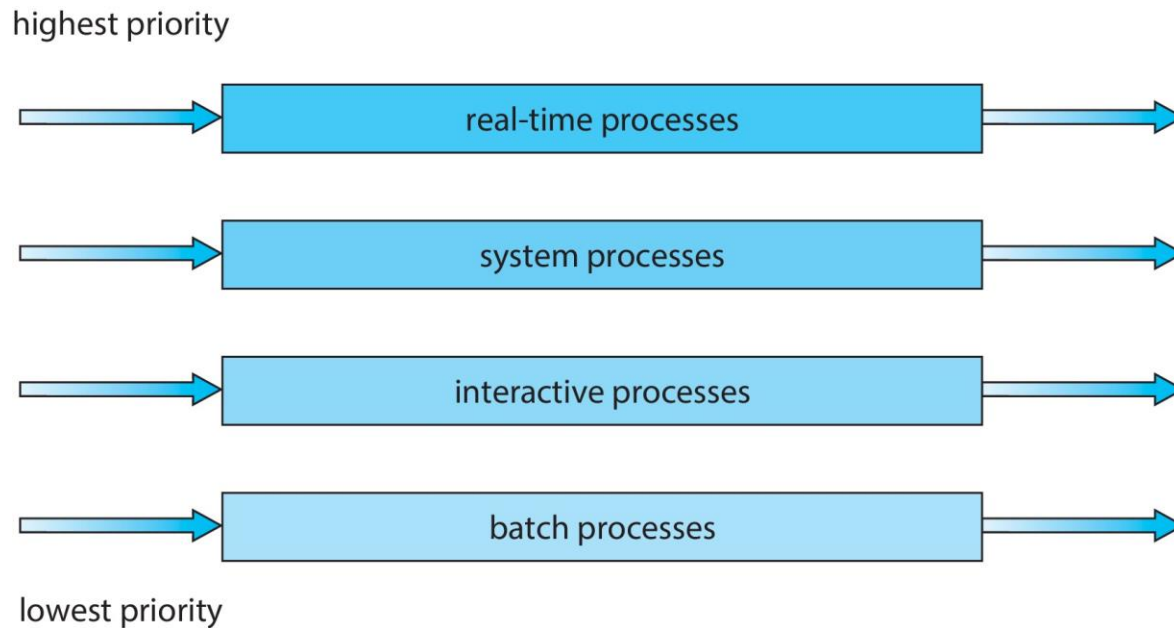
- With priority scheduling, have separate queues for each priority.
- Schedule the process in the highest-priority queue!






Multilevel Queue

- ❑ Prioritization based upon process type
- ❑ Each queue gets certain amount of CPU time (60%, 20%, 10%, 10%)





Multilevel Feedback Queue (MLFQ)

- ❑ A process can move between the various queues; aging can be implemented this way
- ❑ Multilevel-feedback-queue scheduler defined by the following parameters:
 - ❑ number of queues
 - ❑ scheduling algorithms for each queue
 - ❑ method used to determine when to upgrade a process
 - ❑ method used to determine when to demote a process
 - ❑ method used to determine which queue a process will enter when that process needs service






Example of Multilevel Feedback Queue

Three queues:

- Q_0 – RR with time quantum 8 milliseconds
- Q_1 – RR time quantum 16 milliseconds
- Q_2 – FCFS

Scheduling

- A new job enters queue Q_0 which is served FCFS
 - ▶ When it gains CPU, job receives 8 milliseconds
 - ▶ If it does not finish in 8 milliseconds, job is moved to queue Q_1
- At Q_1 job is again served FCFS and receives 16 additional milliseconds
 - ▶ If it still does not complete, it is preempted and moved to queue Q_2




• Result approximates SRTF:

- CPU bound jobs drop like a rock
- Short-running I/O bound jobs stay near top







MLFQ Example



Process	Arrival Time (ms)	Burst Time (ms)
P1	0	10
P2	2	15
P3	5	2
P4	12	14
P5	18	6





Multilevel Feedback Queue (MLFQ)


- ❑ MLFQ is commonly used in many systems such as BSD Unix, Solaris, Window NT and subsequent Window operating systems
- ❑ MLFQ has several distinctive advantages:
 - ❑ It does not require prior knowledge on CPU burst time
 - ❑ It handles interactive jobs by delivering similar performance as that of SJF or SRTF
 - ❑ It is also “fair” by making progress on CPU-bound jobs
- ❑ The possible starvation problem can be handled by reshuffling the jobs to different queues periodically
 - ❑ After some period, move all jobs to the top queue





Thread Scheduling

- When the OS supports threads, the kernel-level threads are the ones being scheduled, not processes, User-level threads are managed by a thread library instead
- The OS uses an intermediate data structure between user and kernel threads, referred as a lightweight process (LWP)
 - It appears to be a virtual processor on which user threads are scheduled to “run”
 - Each LWP attached to kernel thread (one-to-one)
- Under many-to-one and many-to-many models, thread library schedules user-level threads to run on LWP/ This is known as **process-contention scope (PCS)** since scheduling competition takes place among the threads belonging to the same process, typically done via priority set by programmers
- Kernel thread scheduled onto available CPU is **system-contention scope (SCS)** – competition among all threads in the system – CPU scheduling
- Systems using one-to-one mapping model, such as Windows, Linux, and Solaris, schedule threads using only SCS





Multiple-Processor Scheduling


- ❑ CPU scheduling is more complex when multiple CPUs are available – load sharing
- ❑ **Asymmetric multiprocessing** – only one processor accesses the system data structures, alleviating the need for data sharing. The other processors execute only user code
- ❑ **Symmetric multiprocessing (SMP)** – each processor is self-scheduling, all processes in common ready queue, or each has its own private queue of ready processes
 - ❑ All modern OS support SMP, including Window, Linux, and Mac OS X
- ❑ **Processor affinity** – process has affinity for processor on which it is currently running,
 - ❑ When a thread has been running on one processor, the cache content of that processor stores the memory accesses by that thread. We refer to this as a thread having affinity for a processor (i.e. “processor affinity”)
 - ❑ There is a high cost of invalidating and repopulating caches, most SMP systems try to avoid migration of processes from one processor to another
 - ❑ **Soft affinity** – the OS attempt to keep a process running on the same processor, not guaranteeing it
 - ❑ **Hard affinity** – allow a process to specify a subset of processors on which it may run





NUMA and CPU Scheduling

If the operating system is **NUMA-aware**, it will assign memory close to the CPU the thread is running on.





Multiple-Processor Scheduling – Load Balancing



- ❑ On SMP systems, need to keep all CPUs loaded for efficiency
- ❑ **Load balancing** attempts to keep workload evenly distributed
- ❑ **Push migration** – a specific task periodically checks the load on each processor, and if it finds an imbalance, pushes task from overloaded CPU to idle or less-busy CPUs
- ❑ **Pull migration** – idle processors pulls waiting task from a busy processor
- ❑ Push and pull migration need not to be mutually exclusive and are in fact often both implemented in parallel on load-balancing systems
- ❑ Load balancing often counteracts the benefits of processor affinity





Multicore Processors

- Recent trend to place multiple processor cores on same physical chip
 - Faster and consumes less power
- Multiple threads per core also growing
 - Memory stall**: a situation when a processor accesses memory, it spends a significant amount of time waiting for the data to become available, due to various reasons such as cache miss
- The scheduling can take advantage of **memory stall** to make progress on another thread while memory retrieve happens
 - If one thread stalls while waiting for memory, the core can switch to another thread. This becomes a **dual-thread processor core**, or resembles two logical processors
 - A **dual-threaded, dual-core system** presents **four** logical processors to the operating system
 - UltraSPARC T3 CPU has 16 cores per chip and 8 hardware threads per core, from OS perspective, this appears to be 128 logical processors

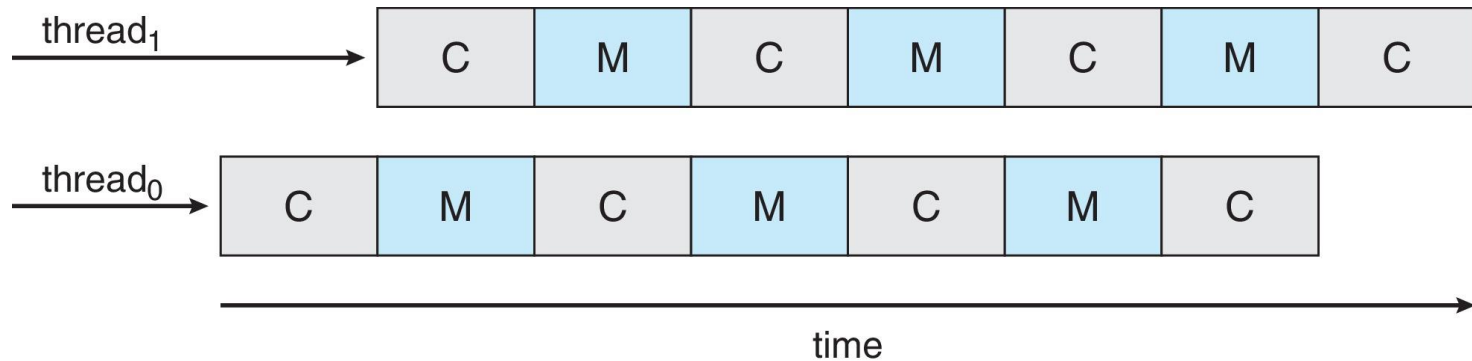




Multithreaded Multicore System

Each core has > 1 hardware threads.




If one thread has a memory stall, switch to another thread!





Multithreaded Multicore System

- ❑ **Chip-multithreading (CMT)** assigns each core multiple hardware threads. (Intel refers to this as **hyperthreading**.)
- ❑ On a quad-core system with 2 hardware threads per core, the operating system sees 8 logical processors.

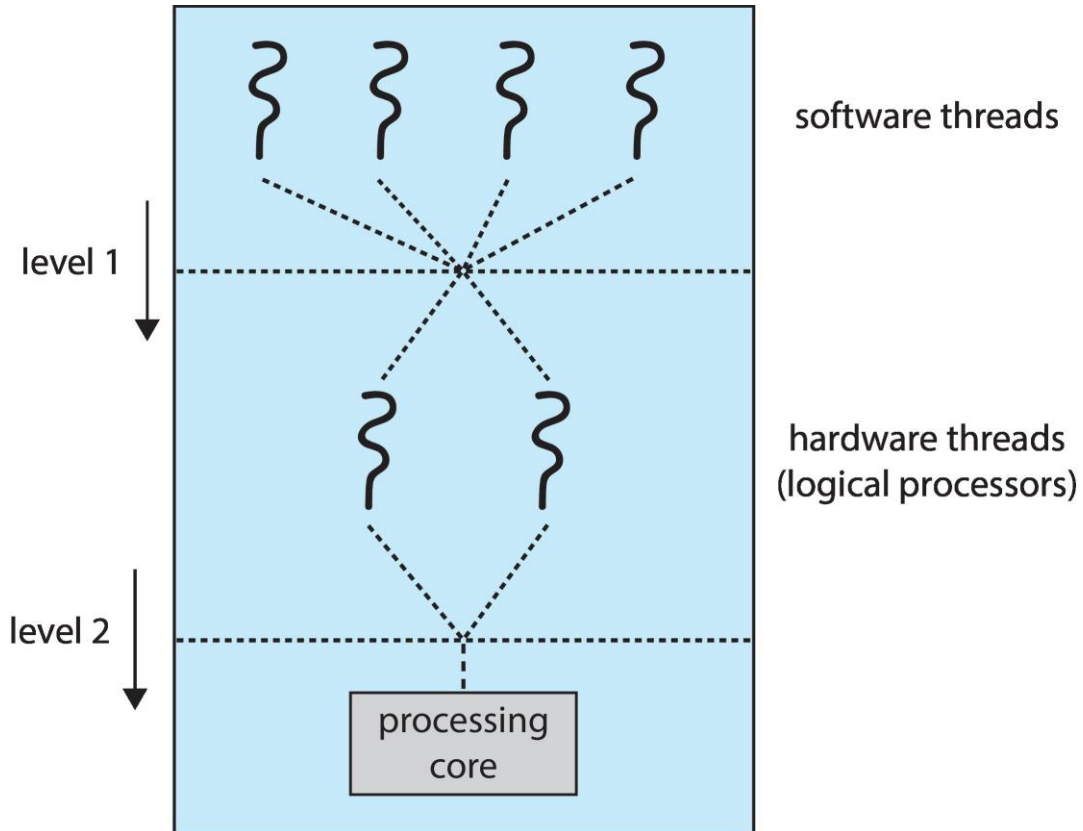




Multithreaded Multicore System

□ Two levels of scheduling:

1. The operating system deciding which software thread to run on a logical CPU
2. How each core decides which hardware thread to run on the physical core.





Real-Time CPU Scheduling


- This requires performance guarantee – predictability
- **Hard real-time systems** – task must be serviced by its deadline
 - Attempt to meet all deadlines
 - EDF - Earlier Deadline First scheduling
- **Soft real-time systems** – Critical real-time tasks have the highest priority, but no guarantee as to when tasks will be scheduled
 - Attempt to meet deadlines with high probability
 - Minimize the miss ratio or maximize completion ratio





Priority-based Scheduling


- For real-time scheduling, scheduler must support preemptive, priority-based scheduling
 - But only guarantees soft real-time
- For hard real-time must also provide ability to meet deadlines
- Processes have new characteristics: **periodic** ones require CPU at constant intervals
 - Has processing time t , deadline d , period p
 - $0 \leq t \leq d \leq p$
 - **Rate** of periodic task is $1/p$





Rate Monotonic Scheduling

- A priority is assigned based on the inverse of its period
 - Shorter periods = higher priority;
 - Longer periods = lower priority
- Suppose P_1 has a period of 50 (also deadline), and processing time 25. P_2 has a period of 80 (also deadline), and processing time 35.
 - Since $50 < 80$, P_1 is assigned a higher priority than P_2
- Process P_2 misses finishing its deadline at time 80






Earliest Deadline First Scheduling (EDF)

- Priorities are assigned according to deadlines:

the earlier the deadline, the higher the priority;

the later the deadline, the lower the priority





Algorithm Evaluation

- How to select CPU-scheduling algorithm for an OS?
- Determine criteria, then evaluate algorithms
- **Deterministic modeling**
 - Takes a particular predetermined workload and defines the performance of each algorithm for that workload
- Consider 5 processes arriving at time 0:

<u>Process</u>	<u>Burst Time</u>
P_1	10
P_2	29
P_3	3
P_4	7
P_5	12






Deterministic Evaluation


- For each algorithm, calculate minimum average waiting time
- Simple and fast, but requires exact numbers for input, applies only to those inputs
- FCS is 28ms:



- Non-preemptive SFJ is 13ms:



- RR is 23ms:





Queueing Models

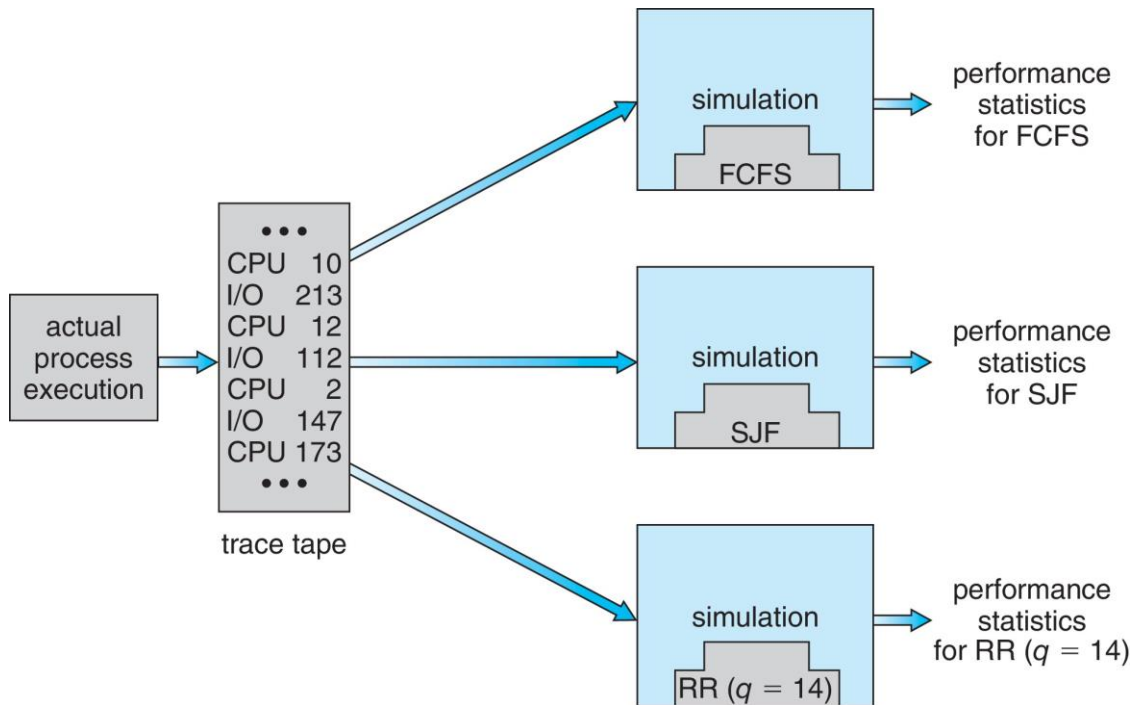
- Mathematical approach for handling stochastic workloads
- Little's Formula
- n = average queue length
- W = average waiting time in queue
- λ = average arrival rate into queue
- Little's law – in steady state, processes leaving queue must equal processes arriving, thus:
$$n = \lambda \times W$$
 - Valid for any scheduling algorithm and arrival distribution
- For example, if on average 7 processes arrive per second, and normally 14 processes in queue, then average wait time per process = 2 seconds





Simulations/Implementations

- ❑ Queueing models are restricted, most with no mathematical solution
- ❑ Simulations or implementation
 - ❑ Build system which allow actual algorithms to run with real data set – more flexible and general



End of Chapter 5

