

Sorting Algorithms

Lecture 7: Priority Queues, Heaps, and Heapsort

Priority Queue: Motivating Example

3 jobs have been submitted to a printer in the order A, B, C.

Consider the printing pool at this moment.

Sizes: Job A — 100 pages

Job B — 10 pages

Job C — 1 page



Average finish time with FIFO service:

$$(100+110+111) / 3 = 107 \text{ time units}$$

Average finish time for shortest-job-first service:

$$(1+11+111) / 3 = 41 \text{ time units}$$

FIFO = First In First Out

Priority Queue: Motivating Example

- The elements in the queue are printing jobs, each with the associated number of pages that serves as its priority
- Processing the shortest job first corresponds to **extracting the smallest element** from the queue
- **Insert** new printing jobs as they arrive

Want a queue capable of supporting two operations:

Insert and **Extract-Min**.

Priority Queue

A *Priority Queue* is an abstract data structure that supports two operations

- **Insert**: inserts the new element into the queue
- **Extract-Min**: removes and returns the smallest element from the queue




Possible Implementations

- unsorted list + a pointer to the smallest element
 - **Insert** in $O(1)$ time
 - **Extract-Min** in $O(n)$ time, since it requires a linear scan to find the new minimum
- sorted (circular) array
 - **Insert** in $O(n)$ time
 - **Extract-Min** in $O(1)$ time
- sorted doubly linked list
 - **Insert** in $O(n)$ time
 - **Extract-Min** in $O(1)$ time *(given pointer to minimum-item)*

Question

Is there any data structure that supports both these priority queue operations in $O(\log n)$ time?


(Binary) Heap




Heaps are “almost complete binary trees”

- All levels are full except possibly the lowest level
- If the lowest level is not full, then nodes must be packed to the left

Heap-order Property



A min-heap



Not a heap


Heap-order property:

The value of a node is at least the value of its parent — **Min-heap**

Heap Properties

- If the heap-order property is maintained, we will show that heaps support the following operations efficiently (n is # elements in the heap):
 - **Insert** in $O(\log n)$ time
 - **Extract-Min** in $O(\log n)$ time
- Structural properties
 - Fact from Discrete Math
A heap of height h has between 2^h to $2^{h+1} - 1$ nodes.
 - If $2^h \leq n < 2^{h+1} \Rightarrow h \leq \log_2 n < h + 1$
 \Rightarrow an n -element heap has height $\Theta(\log n)$.
 - Also, the structure is so regular, it can be represented in an array with no pointers required!!!


Array Implementation of Heap



1	2	3	4	5	6	7	8	9	10
A	B	C	D	E	F	G	H	I	J

- The root is in array position 1
- For any element in array position i
 - The left child is in position $2i$
 - The right child is in position $2i + 1$
 - The parent is in position $\lfloor i/2 \rfloor$

Array Implementation of Heap



1	2	3	4	5	6	7	8	9	10
A	B	C	D	E	F	G	H	I	J

Example: $C = A[3]$.

C's children are


$F = A[2 \cdot 3] = A[6]$ and $G = A[2 \cdot 3 + 1] = A[7]$.

G's parent is $C = A[3] = A[\lfloor 7/2 \rfloor]$.

- The root is in array position 1
- For any element in array position i
 - The left child is in position $2i$
 - The right child is in position $2i + 1$
 - The parent is in position $\lfloor i/2 \rfloor$
- We will draw the heaps as trees, with the understanding that an actual implementation will use simple arrays


Insertion

- Add the new element to the next available position at the lowest level
- Restore the min-heap property if violated
 - General strategy is percolate up (or bubble up): if the parent of the element is larger than the element, then interchange the parent with child.




Insertion

- Add the new element to the next available position at the lowest level
- Restore the min-heap property if violated
 - General strategy is percolate up (or bubble up): if the parent of the element is larger than the element, then interchange the parent with child.




Insertion

- Add the new element to the next available position at the lowest level
- Restore the min-heap property if violated
 - General strategy is percolate up (or bubble up): if the parent of the element is larger than the element, then interchange the parent with child.



Insertion


- Add the new element to the next available position at the lowest level
- Restore the min-heap property if violated
 - General strategy is percolate up (or bubble up): if the parent of the element is larger than the element, then interchange the parent with child.



Percolate up to maintain the min-heap property

Insertion


- Add the new element to the next available position at the lowest level
- Restore the min-heap property if violated
 - General strategy is percolate up (or bubble up): if the parent of the element is larger than the element, then interchange the parent with child.



- Correctness: after each swap, the min-heap property is satisfied for the subtree rooted at the new element
- Time complexity = $O(\text{height}) = O(\log n)$


Insertion (2nd Example)

- Add the new element to the next available position at the lowest level
- Restore the min-heap property if violated
 - General strategy is percolate up (or bubble up): if the parent of the element is larger than the element, then interchange the parent with child.




Insertion (2nd Example)

- Add the new element to the next available position at the lowest level
- Restore the min-heap property if violated
 - General strategy is percolate up (or bubble up): if the parent of the element is larger than the element, then interchange the parent with child.




Insertion (2nd Example)

- Add the new element to the next available position at the lowest level
- Restore the min-heap property if violated
 - General strategy is percolate up (or bubble up): if the parent of the element is larger than the element, then interchange the parent with child.



Insertion (2nd Example)

- Add the new element to the next available position at the lowest level
- Restore the min-heap property if violated
 - General strategy is percolate up (or bubble up): if the parent of the element is larger than the element, then interchange the parent with child.



Percolate up to maintain the min-heap property


In this example, swapping stopped BEFORE reaching the top.

Insertion

Insert(x, i): Add item x to heap $A[1 \cdots i - 1]$ creating heap $A[1 \cdots i]$

```
begin
   $A[i] = x;$ 
   $j = i;$ 
  while  $A[j] < A[\lfloor \frac{j}{2} \rfloor]$  and  $j > 1$  do
    //  $A[j]$  is less than its parent
    Swap  $A[j]$  and  $A[\lfloor \frac{j}{2} \rfloor]$ ; // Bubble Up
     $j = \lfloor \frac{j}{2} \rfloor$ 
  end
end
```


Extract-Min: First Attempt



Min-heap property preserved, **but completeness** not preserved!


Extract-Min

- Copy the last element X to the root (i.e., overwrite the minimum element stored there)
- Restore the min-heap property by percolating (or bubbling down): if the element is larger than either of its children, then interchange it with the smaller of its children.



Extract-Min


- Copy the last element X to the root (i.e., overwrite the minimum element stored there)
- Restore the min-heap property by percolating (or bubbling down): if the element is larger than either of its children, then interchange it with the smaller of its children.



Remove last element and copy its value to the root

Extract-Min


- Copy the last element X to the root (i.e., overwrite the minimum element stored there)
- Restore the min-heap property by percolating (or bubbling) down): if the element is larger than either of its children, then interchange it with the smaller of its children.



Percolate down to maintain min-heap property

Extract-Min


- Copy the last element X to the root (i.e., overwrite the minimum element stored there)
- Restore the min-heap property by percolating (or bubbling down): if the element is larger than either of its children, then interchange it with the smaller of its children.



Percolate down to maintain min-heap property

Extract-Min

- Copy the last element X to the root (i.e., overwrite the minimum element stored there)
- Restore the min-heap property by percolating (or bubbling down): if the element is larger than either of its children, then interchange it with the smaller of its children.



Percolate down to maintain min-heap property

- Correctness: after each swap, the min-heap property is satisfied for all nodes except the node containing X (with respect to its children)
- Time complexity = $O(\text{height}) = O(\log n)$

Extract-Min

Extract-Min(i): Remove (smallest) item $A[1]$ in Heap and make $A[1 \cdots i - 1]$ a Heap of remaining elements. Empty array cells will contain an ∞ as an empty flag.

```
begin
  Output( $A[1]$ );
  Swap  $A[1]$  and  $A[i]$ ;  $A[i] = \infty$ ;  $j = 1$ ; // Remove smallest
   $l = A[2j]$ ;  $r = A[2j + 1]$ ;
  while  $A[j] > \min(l, r)$  do
    // if  $A[j]$  larger than a child, swap with min child
    if  $l < r$  then
      | Swap  $A[j]$  with  $A[2j]$ ;  $j = 2j$ ;
    else
      | Swap  $A[j]$  with  $A[2j + 1]$ ;  $j = 2j + 1$ ;
    end
     $l = A[2j]$ ;  $r = A[2j + 1]$ ;
  end
end
```

Heapsort

- **Build a binary heap of n elements**
 - the minimum element is at the top of the heap
 - insert n elements one by one $\Rightarrow O(n \log n)$
(A more clever approach can do this in $O(n)$ time.)
- **Perform n Extract-Min operations**
 - the elements are extracted in sorted order
 - each Extract-Min operation takes $O(\log n)$ time
 $\Rightarrow O(n \log n)$
- Total time complexity: $O(n \log n)$

Summary

- A *Priority queue* is an abstract data structure that supports two operations: **Insert** and **Extract-Min**.
- If priority queues are implemented using heaps, then these two operations are supported in $O(\log n)$ time.
- Heapsort takes $O(n \log n)$ time, which is as efficient as merge sort and quicksort.

New Operation

Sometimes priority queues need to support another operation called **Decrease-Key**

- **Decrease-Key**: decreases the value of one specified element
- **Decrease-Key** is used in later algorithms, e.g., in Dijkstra's algorithm for finding Shortest Path Trees

Question

How can heaps be modified to support **Decrease-Key** in $O(\log n)$ time?

Going Further

- For some algorithms, there are other desirable Priority Queue operations, e.g., *Delete* an arbitrary item and *Melding* or taking the union of two priority queues
- There is a tradeoff between the costs of the various operations. Depending upon where the data structure is used, different priority queues might be better.
- Most famous variants are *Binomial Heaps* and *Fibonacci Heaps*