

# Lecture 9: Greedy Algorithms

---

Version of February 27, 2019

A greedy algorithm always makes the choice that looks best at the moment and adds it to the current partial solution.

Greedy algorithms don't always yield optimal solutions, but when they do, they're usually the simplest and most efficient algorithms available.

# Outline

1. Interval Scheduling


2. Knapsack

3. Interval Partitioning

# Interval Scheduling

## Interval scheduling.


- Job  $j$  starts at  $s_j$  and finishes at  $f_j$ .
- Two jobs are **compatible** if they don't overlap.
- Goal: find maximum size subset of mutually compatible jobs.



# Interval Scheduling

## Interval scheduling.


- Job  $j$  starts at  $s_j$  and finishes at  $f_j$ .
- Two jobs are **compatible** if they don't overlap.
- Goal: find maximum size subset of mutually compatible jobs.



# Interval Scheduling

## Interval scheduling.

- Job  $j$  starts at  $s_j$  and finishes at  $f_j$ .
- Two jobs are **compatible** if they don't overlap.
- Goal: find maximum size subset of mutually compatible jobs.



# Interval Scheduling: Greedy Algorithms

Greedy template. Consider jobs in some order.

Take a job provided it's *compatible* with the ones already taken.

Usually, many compatible jobs can co-exist.

Need to choose a rule specifying which job to choose next.

Three *possible* rules are:

- [Earliest start time]

Consider jobs in increasing order of start time  $s_j$ .

- [Shortest interval]

Consider jobs in increasing order of interval length  $f_j - s_j$ .

- [Fewest conflicts]

For each job, count the number of conflicting jobs  $c_j$ .

Schedule in ascending order of conflicts  $c_j$ .

# Interval Scheduling: Greedy Algorithms

Greedy template. Consider jobs in some order.

Take a job provided it's compatible with the ones already taken.




Order on earliest start time

Chooses {e} instead of {a,b,c,d}



Order on shortest interval

Chooses {c} instead of {a,b}




Order on fewest conflicts

Chooses {f} which forces choosing {a,f,d} instead of {a,b,c,d}

Examples above provide *counterexamples* for the three rules proposed. For each of them, provides an example input for which following that rule yields a non-optimal (i.e., non max-size) schedule.

## Interval Scheduling: Greedy Algorithm

*Greedy algorithm.* Consider jobs in increasing order of **finish time**. Take each job provided it's compatible with the ones already taken





# Interval Scheduling: Greedy Algorithm

**Greedy algorithm.** Consider jobs in increasing order of **finish time**. Take each job provided it's compatible with the ones already taken.

**Intuition:** leaves maximum interval for scheduling the rest of the jobs.

```
Sort jobs by finish times so that  $f_1 \leq f_2 \leq \dots \leq f_n$   
 $A \leftarrow \emptyset, \text{ last} \leftarrow 0$   
for  $j \leftarrow 1$  to  $n$   
    if  $s_j \geq \text{last}$  then  $A \leftarrow A \cup \{j\}, \text{ last} \leftarrow f_j$   
return  $A$ 
```

**Running time dominated by cost of sorting:**  $\Theta(n \log n)$ .

- Remember the finish time of the last job added to  $A$ .
- Job  $j$  is compatible with  $A$  if  $s_j \geq \text{last}$ .

**Remember:**

Correctness (optimality) of greedy algorithms is usually not obvious.

**Need to prove!**

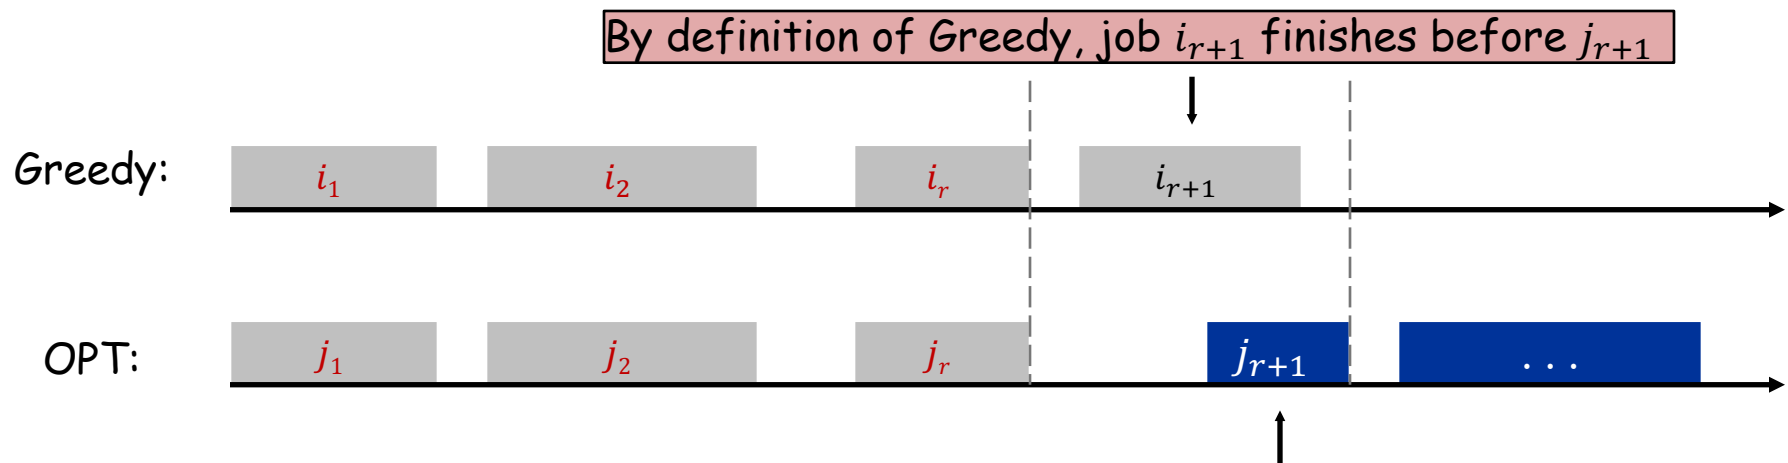
# Interval Scheduling: Correctness

**Theorem.** Greedy algorithm is optimal.

**Proof.**

- Assume Greedy is different from OPT. Let's see what's different.
- Let  $i_1, i_2, \dots, i_k$  denote the set of jobs selected by greedy.
- Let  $j_1, j_2, \dots, j_m$  denote set of jobs in the optimal solution.

Find largest possible value of  $r$  such that  $i_1 = j_1, i_2 = j_2, \dots, i_r = j_r$

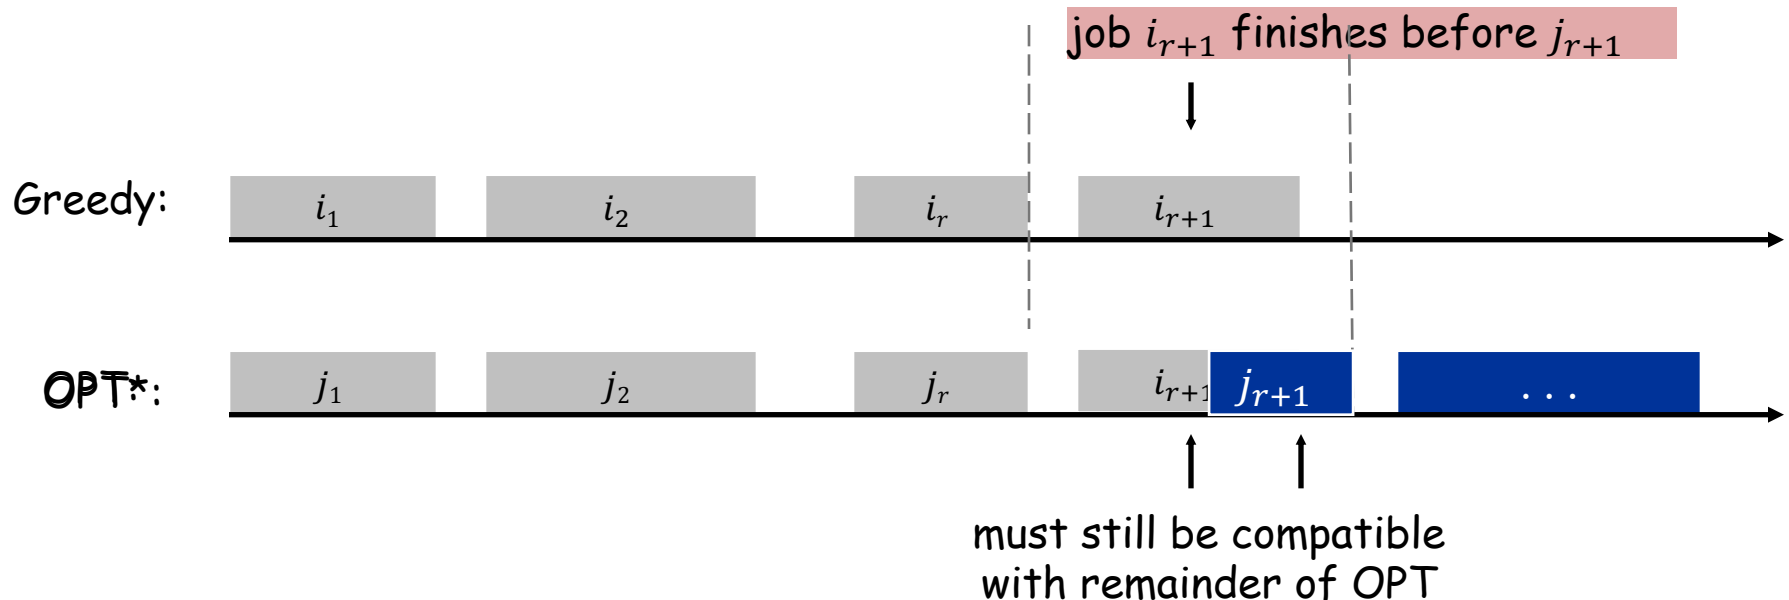


# Interval Scheduling: Correctness

**Theorem.** Greedy algorithm is optimal.

## Proof (Continued)


- Assume Greedy is different from OPT.
- Choose largest  $r$  such that  $i_t = j_t$  for  $t \leq r$  and  $i_{r+1} \neq j_{r+1}$ .
- Create OPT\* from OPT by replacing  $j_{r+1}$  with  $i_{r+1}$ .
- OPT\* is still a legal solution and has same size as OPT.
- $\Rightarrow$  OPT\* is also Optimal



# Interval Scheduling: Correctness

Proof. So far

- Assumed Greedy  $\neq$  OPT.
- Let  $r$  be such that OPT shares first  $r$  items with Greedy.
- Process creates new Optimal solution OPT\*, which shares first  $r+1$  items with Greedy



Can repeat this process starting with Greedy and (optimal) OPT\*  
Continue repeating this process until OPT is the same as greedy.

**Important:** Since cost remains the same, final solution we've created,  
which is Greedy, is optimal!

Finished!

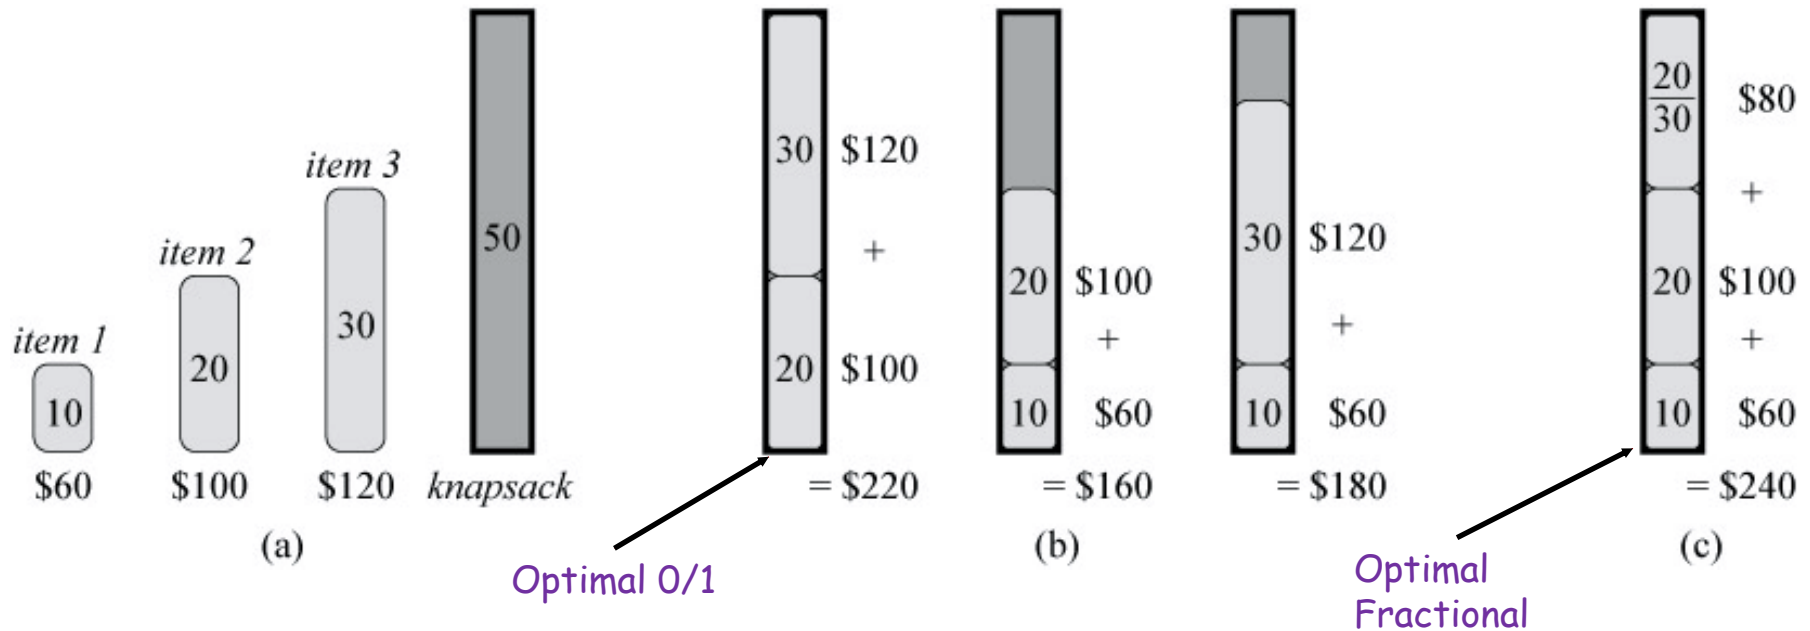
# Outline

1. Interval Scheduling

2. Knapsack

3. Interval Partitioning

# The Fractional Knapsack Problem



**Input:** Set of  $n$  items: item  $i$  has weight  $w_i$  and value  $v_i$ , and a knapsack with capacity  $W$ .

**Goal:** Find  $0 \leq x_1, \dots, x_n \leq 1$  such that  $\sum_{i=1}^n x_i w_i \leq W$  and  $\sum_{i=1}^n x_i v_i$  is maximized.

There are two different versions of this problem:

- The  $x_i$ 's must be 0 or 1: The 0/1 knapsack problem.
- The  $x_i$ 's can take fractional values: The fractional knapsack problem

# The Greedy Algorithm for Fractional Knapsack

```
Sort items so that  $\frac{v_1}{w_1} > \frac{v_2}{w_2} > \dots > \frac{v_n}{w_n}$   
 $w \leftarrow W$   
for  $i \leftarrow 1$  to  $n$   
    if  $w_i \leq w$  then  
         $x_i \leftarrow 1$   
         $w \leftarrow w - w_i$   
    else  
         $x_i \leftarrow w/w_i$   
    return  
return
```

## Idea:

- Sort all items by value-per-pound
- For each item, take as much as possible

Running time:  $\Theta(n \log n)$

Note: This algorithm cannot solve the 0/1 version optimally.

## Greedy Algorithm: Correctness

**Theorem:** The greedy algorithm is optimal.

**Proof:** We assume that  $\sum_{i=1}^n w_i \geq W$ , so knapsack is fully packed. Otherwise the algorithm is trivially optimal.

Let the greedy solution be  $G = (x_1, x_2, \dots, x_k, 0, \dots, 0)$

- Note: for  $i < k$ ,  $x_i = 1$ ; for  $i > k$ ,  $x_i = 0$ ;  $0 \leq x_k \leq 1$ .

Consider any optimal solution  $O = (y_1, y_2, \dots, y_n)$

- Note: Since both  $G$  and  $O$  must fully pack the knapsack,

$$\sum_{i=1}^k x_i w_i = W = \sum_{i=1}^n y_i w_i$$

Look at the first item  $i$  where the two solutions differ.

$$G = x_1 \ x_2 \ x_3 \ \dots \ x_{i-1} \ \textcolor{red}{x_i} \ \dots \ x_k \ \dots \ 0 \ \dots \ 0$$

$$O = x_1 \ x_2 \ x_3 \ \dots \ x_{i-1} \ \textcolor{red}{y_i} \ \dots \ y_k \ \dots \ y_{n-1} \ y_n$$

By definition of greedy,  $x_i > y_i$

- Let  $x = x_i - y_i > 0$



## Greedy Algorithm: Correctness (continued)

We now modify  $O$  as follows:

- Set  $y_i \leftarrow x_i$  and remove amount of some items in  $i + 1$  to  $n$  of total weight  $x_i w_i$
- This is always doable because in both  $O$  and  $G$ , the used total weight of items  $i$  to  $n$  is the same. ( $\sum_{i=1}^k x_i w_i = W = \sum_{i=1}^n y_i w_i$ )

After the modification:

- The total value of this new  $O$  has not decreased, since all the items  $i + 1$  to  $n$  have lesser or equal value-per-pound than item  $i$
- This new  $O$ 's value can not be greater than before, since  $O$  was already an optimal solution,
- $\Rightarrow O$ 's value stays the same  $\Rightarrow O$  is still an optimal solution
- $O$ 's first index that differs from  $G$  is now at least  $i + 1$

Repeating this process will eventually convert  $O$  into  $G$

(1<sup>st</sup> index that differs always increases so the process must stop).

Note that process keeps solution value of  $O$  invariant (and optimal).

$\Rightarrow$  at end  $G = O \Rightarrow$  Greedy is optimal

# Outline

1. Interval Scheduling

2. Knapsack


3. Interval Partitioning

# Interval Partitioning

## Interval partitioning.

- Lecture  $j$  starts at  $s_j$  and finishes at  $f_j$ .
- Goal: find the minimum number of classrooms to schedule all lectures so that no two occur at the same time in the same room.

Ex: This schedule uses 4 classrooms to schedule 10 lectures.




# Interval Partitioning

Interval partitioning.

- Lecture  $j$  starts at  $s_j$  and finishes at  $f_j$ .
- Goal: find the minimum number of classrooms to schedule all lectures so that no two occur at the same time in the same room.

Ex: This schedule uses only 3.



## Interval Partitioning: Greedy Algorithm

**Greedy algorithm.** Consider lectures in increasing order of start time: assign lecture to any compatible classroom.

```
Sort intervals by starting time so that  $s_1 \leq s_2 \leq \dots \leq s_n$ .  
  
 $d \leftarrow 0$  // # classrooms used so far  
  
for  $j \leftarrow 1$  to  $n$   
    if lecture  $j$  is compatible with some classroom  $k$  then  
        schedule lecture  $j$  in classroom  $k$   
    else  
        allocate a new classroom  $d + 1$   
        schedule lecture  $j$  in classroom  $d + 1$   
         $d \leftarrow d + 1$ 
```

Greedy! It only opens a new classroom if it is needed.

Need to prove optimality.

More specifically, need to show that processing the lectures **ordered by starting time** implies optimality.


*Convince yourself that processing by finishing time does NOT yield optimal solution!*

# Interval Partitioning

## Interval partitioning.

- Lecture  $j$  starts at  $s_j$  and finishes at  $f_j$ .
- Goal: find the minimum number of classrooms to schedule all lectures so that no two occur at the same time in the same room.

ALG: Sort by start time. Insert in order, opening new classroom when needed



# Interval Partitioning: Lower Bound on Optimal Solution


**Def.** The **depth** of a set of open intervals is the maximum number that exist at any instant of time.

*Intuition:* At each second of the day keep track of how many simultaneous classes are being taught at that time. The depth is the maximum number you have seen.

**Key observation.** Minimum number of classrooms needed  $\geq$  **depth**.

**Ex:** Depth of schedule below = 3  $\Rightarrow$  this schedule is optimal.

**We will show:** The # classrooms used by the greedy algorithm = **depth**.



## Interval Partitioning: Correctness

**Theorem.** Greedy algorithm is optimal.

**Pf.**

- Let  $d$  = number of classrooms opened by greedy algorithm .
- Classroom  $d$  is opened because we needed to schedule a lecture, say  $j$ , that is incompatible with all  $d - 1$  other classrooms.
- Since **we sorted by start time**, all these incompatibilities are caused by lectures that all start before (or at same time as)  $s_j$  and finish later than  $s_j$ .
- Thus, we have  $d$  lectures all overlapping at time  $s_j + \varepsilon$  for some  $\varepsilon > 0$ ; the  $d-1$  incompatible ones and lecture  $j$ ,
- $\Rightarrow$  **depth**  $\geq d$ .
- Minimum number of classrooms needed  $\geq$  **depth**  $\geq d$ ,  
Since Greedy uses only  $d$ , Greedy is optimal.



## Interval Partitioning: Running Time

```
Sort intervals by starting time so that  $s_1 \leq s_2 \leq \dots \leq s_n$ .  
 $d \leftarrow 0$  // # classrooms used so far  
for  $j \leftarrow 1$  to  $n$   
    if lecture  $j$  is compatible with some classroom  $k$  then (*)  
        schedule lecture  $j$  in classroom  $k$  (**)  
    else  
        allocate a new classroom  $d + 1$   
        schedule lecture  $j$  in classroom  $d + 1$  (***)  
         $d \leftarrow d + 1$ 
```

- To implement line (\*) the algorithm maintains, for each classroom, the finishing time of the last item placed in the classroom. It then compares  $s_j$  to those finishing times. If  $s_j \geq$  one of those finishing times, it places lecture  $j$  in the associated classroom
- A Brute-force implementation of line (\*) uses  $O(n)$  time  
 $\Rightarrow O(n^2)$  in total
- Observation: If  $j$  is not compatible with the classroom with **the earliest finish time**, then  $j$  is not compatible with any other classroom

## Interval Partitioning: Running Time

```
Sort intervals by starting time so that  $s_1 \leq s_2 \leq \dots \leq s_n$ .
 $d \leftarrow 0$  // # classrooms used so far
for  $j \leftarrow 1$  to  $n$ 
    if lecture  $j$  is compatible with some classroom  $k$  then (*)
        schedule lecture  $j$  in classroom  $k$  (**)
    else
        allocate a new classroom  $d + 1$ 
        schedule lecture  $j$  in classroom  $d + 1$  (***)
         $d \leftarrow d + 1$ 
```

Running time:  $\Theta(n \log n)$

- To implement line (\*) we can keep the classrooms in a **min priority queue** using the finishing times of the last class in the room as key

$O(\log n)$  ▪ To check whether there is a compatible classroom we do an **extract-min** to find min finishing time  $f_{\min}$  in the priority queue and check whether  $f_{\min} \leq s_j$

$O(\log n)$  ▪ To implement (\*\*) just add the new finishing time  $f_j$  to p. queue

$O(\log n)$  ▪ To implement (\*\*\*) re-insert  $f_{\min}$  back into p. queue AND insert the new finishing time  $f_j$  into p. queue